

# Multi Objective PPO

[git@github.com:Yaishnavi/MultiObjective-PPO.git](https://github.com/Yaishnavi/MultiObjective-PPO)

## Contents

Introduction  
Datasets  
Experiment  
Results  
Installation  
References

## Introduction

Multi-Objective RL problem with  $k$  different reward functions. Find a policy that performs well for all  $k$  reward functions.

Implemented the following algorithm:

Use PPO to find policy  $\pi_1$  with respect to rewards  $r_1$

For  $k = 2$  to  $K$

    Use PPO to find policy  $\pi_k$  where reward is  $r_k + \alpha \text{KL}(\pi_{k-1} || \pi_k)$

The **Kullback-Leibler Divergence** score, or KL divergence score, quantifies how much one probability distribution differs from another probability distribution.

The KL divergence between two distributions  $Q$  and  $P$  is often stated using the following notation:  $\text{KL}(P || Q)$

Where the “ $||$ ” operator indicates “divergence” or  $P$ s divergence from  $Q$ .

**PPO** is a policy gradient method for reinforcement learning, which alternates between sampling data through interaction with the environment, and optimizing a "surrogate" objective function using stochastic gradient ascent. Whereas standard policy gradient methods perform one gradient update per data sample, we propose a novel objective function that enables multiple epochs of minibatch updates.

## Proximal Policy Optimization

PPO can be viewed as an approximation of TRPO, but unlike TRPO, which uses a second-order Taylor expansion, PPO uses only a first-order approximation, which makes PPO very effective in RNN networks and in a wide distribution space.

```

for  $i \in 1, 2, \dots, N$  do:
  Run policy  $\pi_\theta$  for  $T$  timesteps, collecting  $s_t, a_t, r_t$ 
   $\pi_{\text{old}} \leftarrow \pi_\theta$ 
  for  $j \in 1, 2, \dots, M$  do:
    
$$J_{\text{PPO}}(\theta) = \sum_{t=1}^T \frac{\pi_\theta(a_t|s_t)}{\pi_{\text{old}}(a_t|s_t)} \hat{A}_t - \lambda \text{KL}[\pi_{\text{old}}|\pi_\theta]$$

    Update  $\theta$  by a gradient method w.r.t.  $J_{\text{PPO}}(\theta)$ 
  end for
  for  $j \in 1, 2, \dots, B$  do:
    
$$L_{\text{BL}}(\phi) = - \sum_{t=1}^T \left( \sum_{t'>t} \gamma^{t'-t} r_{t'} - V_\phi(s_t) \right)^2$$

    Update  $\phi$  by a gradient method w.r.t.  $L_{\text{BL}}(\phi)$ 
  end for
  if  $\text{KL}[\pi_{\text{old}}|\pi_\theta] > \beta_{\text{high}} \text{KL}_{\text{target}}$  then
     $\lambda \leftarrow \alpha \lambda$ 
  if  $\text{KL}[\pi_{\text{old}}|\pi_\theta] < \beta_{\text{high}} \text{KL}_{\text{target}}$  then
     $\lambda \leftarrow \alpha / \lambda$ 

```

The first half of Estimate Advantage is obtained through the rollout strategy, and the second half of  $V$  is obtained from a value network. (Value network can be trained by the data obtained by rollout, where the mean square error is used).

There is a clipped surrogate objective,

$$L^{\text{CLIP}}(\theta) = \hat{E}_t[r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t]$$

where  $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$

### Dataset

Fruit API is a universal deep reinforcement learning framework, which is designed meticulously to provide a friendly user interface, a fast algorithm prototyping tool, and a multi-purpose library for the RL research community. External environments can be integrated into the framework easily by plugging into FruitEnvironment. Finally, we developed 5 extra environments as a testbed to examine different disciplines in deep RL:

- Mountain car (multi-objective environment/graphical support)
- Deep sea treasure (multi-objective environment/graphical support)

- Tank battle (multi-agent/multi-objective/human-agent cooperation environment)
- Food collector (multi-objective environment)
- Other environments used but not evaluated for this task are Milk factory and fruit tree.


	Number of rewards/objectives	Number of actions	State Dimension
Mountain Car	3	3	1
Deep Sea Treasure	2	4	2
Tank Battle	2	5	650x650x3
Food Collector	2	6	300x325x3
Mountain Car 2	2	3	1

### Experiment


PPO algorithm is used to compute the policy for each objective/reward function. This objective is updated with the KL divergence of the previous policy. The baseline using a linear combination of the rewards to determine the best policy. Here dense neural network model is used for all of the tasks mentioned above.

### Results

#### 1. Mountain Car




Proposed method




Baseline

## 2. Deep Sea Treasure




Proposed method




baseline

### 3. Tank Battle




Proposed Method




baseline

#### 4. Food Collector




Proposed method




baseline

## 5. Mountain Car2



Proposed Method



baseline

## **Installation**

1. Install fruit api to access environments  
Git clone <https://github.com/garlicdevs/Fruit-API.git>  
Open install path and run python setup.py install
2. Other packages needed: Python==3.6, Pytorch, Matplotlib, Numpy

## **References**

- Dataset: FruitAPI <https://fruitlab.org/>
- Generalizing across Multi-Objective Reward Functions in Deep Reinforcement Learning <https://arxiv.org/pdf/1809.06364.pdf>
- Proximal Policy Optimization Algorithms <https://arxiv.org/pdf/1707.06347.pdf>
- Multi Objective Deep Reinforcement Learning <https://arxiv.org/pdf/1610.02707.pdf>